

2021 Fall AI701 Homework assignment #2

Read the following problem sheet and answer to the questions. Along with the completed version of this ipynb file, **you should also submit an additional file including solutions to the problems marked with ***.

Problem 1 (20 pts)

(a) (3 pts) Let $X \sim \text{Exp}(\lambda)$ be an exponential random variable with parameter $\lambda > 0$ with probability density function (PDF)

$$f_X(x) = \lambda e^{-\lambda x} \mathbf{1}[x \geq 0].$$

Using the function to sample from a uniform distribution `npr.rand`, implement the following code to sample X .

```
In [ ]: import numpy as np
import numpy.random as npr

# Draw n samples from an exponential distribution with parameter lamb.
def rand_exp(lamb, n):
    #####
    # put your code here
    #####
```

★ (b) (2 pts) We will use the law of large numbers to verify your code. For that, compute the mean and variance of X with PDF

$$f_X(x) = \lambda e^{-\lambda x} \mathbf{1}[x \geq 0].$$

(c) (3 pts) Complete the following code and show that your implementation for `rand_exp` is correct.

```
In [ ]: # compute the mean of an exponential distribution with parameter lamb.
def exp_mean(lamb):
    #####
    # put your code here
    #####

# compute the variance of an exponential distribution with parameter lamb.
def exp_var(lamb):
    #####
    # put your code here
    #####

lamb = 1.5
n = 100000
x = rand_exp(lamb, n)

print(f'true mean {exp_mean(lamb)}, empirical mean {x.mean()}')
print(f'true variance {exp_var(lamb)}, empirical variance {x.var()}')
```

★ (d) (5 pts) Let $Y_1 \sim \text{Gamma}(a_1, b)$ be a gamma random variable with shape parameter $a_1 > 0$ and rate parameter $b > 0$ with PDF

$$f_{Y_1}(y) = \frac{b^{a_1} y^{a_1-1} e^{-by}}{\Gamma(a_1)} \mathbf{1}[y \geq 0].$$

Similarly, let $Y_2 \sim \text{Gamma}(a_2, b)$ be another gamma random variable with shape $a_2 > 0$ and rate $b > 0$. Assuming Y_1 and Y_2 are independent, show that $Y = Y_1 + Y_2 \sim \text{Gamma}(a_1 + a_2, b)$.

★ (e) (4 pts) Explain how to draw samples from a gamma distributed random variable $Y \sim \text{Gamma}(m, 1)$ with $m \in \mathbb{N}$ (\mathbb{N} is a set of natural numbers) using `rand_exp`.

(f) (3 pts) Complete the following codes to sample from $Y \sim \text{Gamma}(m, 1)$ and verify the correctness.

```
In [ ]: # Draw n samples from a gamma random variable with shape m (natural number) and rate 1.
def rand_gamma(m, n):
    #####
    # put your code here
    #####

def gamma_mean(m):
    return m

def gamma_var(m):
    return m

n = 100000
```

```
m = 6
x = rand_gamma(m, n)

print(f'true mean {gamma_mean(m)}, empirical mean {x.mean()}')
print(f'true variance {gamma_var(m)}, empirical variance {x.var()}')
```

Problem 2 (30 pts)

Consider a random variable X whose unnormalized PDF is given as follows:

$$\tilde{f}_X(x) = x^2 \sin(2\pi x) \mathbf{1}[-2 \leq x \leq 2], \quad f_X(x) = \frac{\tilde{f}_X(x)}{Z}, \quad Z := \int_{-2}^2 \tilde{f}_X(x) dx.$$

(Actually, we can compute Z analytically, but for now assume that we can't for some reason). We'd like to simulate X via rejection sampling. For that, we design a proposal distribution with density function $q(x)$ given as

$$q(x) = \frac{3}{16} x^2 \mathbf{1}[-2 \leq x \leq 2],$$

and set $M = 16/3$. Then clearly $Mq(x) \geq \tilde{f}_X(x)$ for all $x \in [-2, 2]$. The following code plotting \tilde{f}_X and q demonstrate this.

```
In [ ]: import matplotlib.pyplot as plt
from scipy.integrate import quad

def f_tilde(x):
    return x**2 * np.sin(2*np.pi*x)**2

def q(x):
    return 3*x**2/16

M = 16.0/3.0

plt.figure()
x = np.linspace(-2, 2, 1000)
plt.plot(x, M*q(x), '-r', label=r'$Mq(x)$')
plt.plot(x, f_tilde(x), '-b', label=r'$\tilde{f}(x)$')
plt.legend()
```

★ (a) (8 pts) Describe how to sample from $q(x)$ using `np.random`, and compute the mean and variance of the distribution.

(b) (4 pts) Complete the following codes to sample from $q(x)$ and verify the correctness.

```
In [ ]: # draw n samples from the distribution with PDF q(x).
def rand_q(n):
    #####
    # put your code here
    #####

q_true_mean = ##### put your answer here #####
q_true_var = ##### put your answer here #####

n = 100000
x = rand_q(n)

print(f'true mean {q_true_mean}, empirical mean {x.mean()}')
print(f'true variance {q_true_var}, empirical variance {x.var()}')
```

(c) (9 pts) Complete the following code to sample from $f_X(x)$ with rejection sampling using $q(x)$ as a proposal. Verify your implementation with the provided codes drawing histograms of samples and comparing empirical means and variances to numerically computed means and expectations.

```
In [ ]: # draw a sample from f_X with rejection sampling.
def rand_f_rejection():
    #####
    # put your code here
    #####

# draw samples
n = 100000
x = np.array([rand_f_rejection() for _ in range(n)])

# verification by drawing empirical distribution

plt.hist(x, bins=100, density=True, facecolor='b', edgecolor='k', alpha=0.5);

# numerical integration for the normalization constant
Z = quad(f_tilde, -2, 2)[0]
```

```
tx = np.linspace(-2, 2, 1000)
plt.plot(tx, f_tilde(tx)/Z, '-r', linewidth=2.0)

# verification by comparing mean and variances
numerical_mean = quad(lambda x: x*f_tilde(x)/Z, -2, 2)[0]
numerical_var = quad(lambda x: x**2*f_tilde(x)/Z, -2, 2)[0] - numerical_mean**2

print(f'numerical mean {numerical_mean}, empirical mean {x.mean()}')
print(f'numerical variance {numerical_var}, empirical variance {x.var()}')
```

(d) (9 pts) Consider the following expectation of X ,

$$E[g(X)] = \int_{-2}^2 g(x) f_X(x) dx = \int_{-2}^2 g(x) \frac{\tilde{f}_X(x)}{Z} dx.$$

Complete the following code to estimating this expectation via self-normalized importance sampling. Use $q(x)$ implemented above as a proposal distribution.

```
In [ ]: # compute the approximation for the expectation E[g(X)] using n number of samples.
# g: target function to compute expectation
# n: number of samples being used
def SNIS(g, n):
    #####
    # put your code here
    #####

# verify your implementation using numerical integration
g = lambda x: np.log(np.abs(x) + 1.0)
n = 100000
snis_expec = SNIS(g, n)

numerical_expec = quad(lambda x: g(x)*f_tilde(x)/Z, -2, 2)[0]

print(f'Expectation via numerical integration {numerical_expec}, expectation via SNIS {snis_expec}')
```

Problem 3 (20 pts)

Let X be a d -dimensional categorical random variable with Probability Mass Function (PMF)

$$f_X(x) = \prod_{j=1}^d \pi_j^{1[x=j]}.$$

Using only `npr.rand`, complete the following code to draw samples from the categorical distribution. Verify the correctness of your implementation using the provide code.

```
In [ ]: # draw a sample from a categorical distribution with parameter pi
def rand_cat(pi):
    #####
    # put your code here
    #####

# verify the code using LLN
d = 5
pi = npr.rand(d)
pi = pi / pi.sum()

n = 100000
x = np.zeros(d)
for i in range(n):
    x[rand_cat(pi)] += 1

print(f'empirical mean {x/x.sum()}')
print(f'true mean {pi}')
```

Problem 4 (30 pts)

Consider a bivariate random variable $X = (X_1, X_2)$ with the following unnormalizing PDF:

$$\tilde{f}_X(x_1, x_2) = \exp\left(-\left(x_1 - \frac{x_2^2}{4}\right)^2 - \frac{x_2^2}{4}\right).$$

You can see a contour plot of the density by the following code.

```
In [ ]: # x: n times 2 matrix or 2 dimensional vector
def log_f_tilde(x):
```

```

if x.ndim == 2:
    x1, x2 = x[:,0], x[:,1]
else:
    x1, x2 = x

return -(x2 - x1**2/4)**2 - x1**2/4

def plot_density(alpha=1.0):
    nx, ny = 50, 50
    x = np.linspace(-5, 5, nx)
    y = np.linspace(-2, 4, ny)
    xx, yy = np.meshgrid(x, y)
    z = np.exp(log_f_tilde(np.concatenate([xx.reshape((-1, 1)), yy.reshape((-1, 1))], -1)))
    plt.contour(x, y, z.reshape((nx, ny)), cmap='inferno', alpha=alpha)

plot_density()

```

We want to write a Metropolis-Hastings algorithm based Markov-chain Monte-Carlo sampler for this distribution. we will use a simple random walk Gaussian distribution as our proposal,

$$q(x' | x) = \mathcal{N}(x' ; x, \sigma^2 I),$$

where we fix $\sigma = 0.2$. Complete the following code to run the random walk Metropolis Hastings. Using the provided code, verify the correctness of your implementation by visualizing samples.

```

In [ ]: # propose a sample from the proposal distribution  $q(x' | x) = \mathcal{N}(x' ; x, \sigma^2 I)$ 
def q_MH(x, sigma):
    #####
    # put your code here
    #####

# run a random-walk Metropolis Hastings
# x0: initial sample
# num_samples: number of samples to collect
# sigma: variance for the proposal distribution
def RW_MH(x0, num_samples, sigma):
    #####
    # put your code here
    #####

# randomly initialize a chain
x0 = 0.01*npr.randn(2)
# collect 10000 samples from RW MH
x = RW_MH(x0, 10000, 0.2)

# visualize the samples
plot_density(alpha=1.0)
plt.scatter(x[:,0], x[:,1], alpha=0.1, color='b')

```

Processing math: 100%